

“Getting Started with the tmp18rt-2022”

This template is a special version of [Tpl18](#) - the current final evolution of a long list of ‘templates’, that started with [EasyCE](#), a minimalistic code base for writing Windows CE games / graphics applications without worrying about OS base code. It evolved in various versions for [IGAD](#), then [UU](#), then IGAD again, and in the meantime it has been used to start virtually all my personal mini-projects. In practice, it is great as a basic starting point, but very limited at the same time.

To use the template:

- you simply extract it from the zip file to a directory of your choice
- you open the .sln file using Visual Studio.

At the time of writing, Visual Studio 2022 Community Edition is an excellent choice. [Get it for free](#), install it using the default options, and you’re good to go.

From here on:

You’re welcome to inspect any file in the project. The interesting content however is in `renderer.cpp`. This file implements a minimal ray tracer, with a stationary camera and a hardcoded scene. In `Renderer::Tick`, the ray tracer loops over lines of pixels:

```
// lines are executed as OpenMP parallel tasks (disabled in DEBUG)
#pragma omp parallel for schedule(dynamic)
for (int y = 0; y < SCRHEIGHT; y++)
{ ... }
```

Like the comment says, lines may be executed in parallel, for better performance. Each line is processed by two loops. The first one does the actual ray tracing, via `Renderer::Trace`:

```
// trace a primary ray for each pixel on the line
for (int x = 0; x < SCRWIDTH; x++)
    accumulator[x + y * SCRWIDTH] =
        float4( Trace( camera.GetPrimaryRay( x, y ) ), 0 );
```

The result of this is a floating point color for each pixel. This needs to be converted to something that can be shown on the screen, which happens in the second loop:

```
// translate accumulator contents to rgb32 pixels
for (int dest = y * SCRWIDTH, x = 0; x < SCRWIDTH; x++)
    screen->pixels[dest + x] =
        RGBF32_to_RGB8( &accumulator[x + y * SCRWIDTH] );
```

From here, you’re on your own. Explore what happens in `Renderer::Trace` (spoiler: not much), and follow the code path into `scene.h`, which you can find in the template folder. The scene is entirely hardcoded, which is great for speed: a high-end PC should get close to 1 billion ray/scene intersections per second. This will come in handy once we get to path tracing.

A few more things about handling pixels in the template:

- A pixel is plotted using `screen->Plot(x, y, color)`.
- The size of the screen can be obtained from `SCRWIDTH` and `SCRHEIGHT`.
- A ‘color’ is a 32-bit unsigned value, where red starts at bit 16, green at 8 and blue at 0. Each color component has a range of 0..255.
- You can write debugging info to the text window using `printf`.



From here: draw what you need using `screen->Plot` and other `Surface` methods, handle keys and mouse input using the methods of the `Renderer` class (see `renderer.h`) and add `.cpp` and `.h` files to extend and structure your project.

Basic math classes can be found in `precomp.h` (starting at line 189). Here you will find `float2`, `float3`, `float4` as well as `int` and `uint` counterparts, with an extensive set of operators. There are also basic classes for storing bounding boxes and for matrix calculations. As with the rest of the template, this serves as a basis; you may find it desirable to add some code of your own depending on what your project needs.

GPGPU

The template provides [OpenCL](#) support to deploy the GPU in your calculations. It's use is demonstrated in the `#if 1 / #endif` block in `myapp.cpp`:

```
static Kernel* kernel = 0;           // statics should be members of MyApp of course.
static Surface bitmap( 512, 512 ); // having them here allows us to disable the OpenCL
static Buffer* clBuffer = 0;         // demonstration using a single #if 0.
if (!kernel)
{
    // initialize OpenCL; compile and load kernel "render" from file "kernels.cl"
    Kernel::InitCL();
    kernel = new Kernel( "cl/kernels.cl", "render" );
    // create an OpenCL buffer over using bitmap.pixels
    clBuffer = new Buffer( 512 * 512, Buffer::DEFAULT, bitmap.pixels );
}
// pass arguments to the OpenCL kernel
kernel->SetArgument( 0, clBuffer );
// run the kernel; use 512 * 512 threads
kernel->Run( 512 * 512 );
// get the results back from GPU to CPU (and thus: into bitmap.pixels)
clBuffer->CopyFromDevice();
// show the result on screen
bitmap.CopyTo( screen, 500, 200 );
```

The code demonstrates the most important steps in writing GPGPU code: loading and compiling a kernel, creating buffers to pass data between 'host' and 'device', setting kernel arguments, executing a kernel on the device, and retrieving data from device to host.

A full OpenCL tutorial is outside the scope of this document. If you want to see an example of OpenCL used in the ADVGRtmp18, please refer to the [voxel template](#) on GitHub.

Note that unlike plain OpenCL, the template allows you to use `#include` files in your OpenCL code. The most common use case for this is the `common.h` file, which gets included by `myapp.cpp` (via `precomp.h`) and the example `kernels.cl` file, to share the default screen resolution between host and device code.

The template also makes some definitions available to your OpenCL code:

- `ISNVIDIA` will be defined if your code is running on NVIDIA hardware;
- `ISAMD` and `ISINTEL` provide the same info, but for AMD and Intel;
- Use `ISAMPERE`, `ISTURING`, `ISPASCAL` to write code specific to an NVIDIA architecture.

Go Forth and Code!

That should do the job for now; if you have any questions do not hesitate to contact me:

bikker.j@gmail.com / j.bikker@uu.nl / bikker.j@buas.nl